Career Guidance Chatbot

API pricing and operating-cost report

Prepared 12 November 2026 | Currency: USD | Deployment assumption: Vercel + Upstash Redis

Purpose. Compare Tavily search costs and LLM API options for a career-guidance chatbot expected to handle variable traffic (approximately 100 to 1,000 user requests per day).

Executive recommendation. Start development with Tavily's free tier and Gemini API's eligible free tier using non-sensitive test data. Measure real search frequency and token usage before upgrading. For production resumes and personal information, review provider data-use terms and choose an appropriate paid/API tier.

Important pricing note

Provider pricing, model availability, free quotas and billing terms change frequently. The figures below reproduce the estimates in the accompanying chat report and are planning figures, not a binding quote. Verify each provider's live pricing page and account dashboard before enabling production billing. Some model names or prices may have changed since publication.

1. Tavily: what counts as a credit?

A Tavily credit is a unit of API consumption. It is not always equal to one user request: the credit cost depends on which endpoint is called and how many URLs are successfully extracted.

Operation	Credit consumption
Basic Search	1 credit per request
Advanced Search	2 credits per request
Basic Extract	1 credit per 5 successful URL extractions
Advanced Extract	2 credits per 5 successful URL extractions

Example — search only: One Basic Search consumes 1 credit.

Example — search plus extraction: One Basic Search followed by Basic Extract on five successfully extracted URLs consumes approximately 2 credits total (1 + 1).

Example — advanced search plus extraction: One Advanced Search plus Basic Extract on five URLs consumes approximately 3 credits (2 + 1). Failed URL extractions are not charged under the documented extraction-credit rules.

A single user question can consume multiple credits if the chatbot performs multiple searches or extracts multiple groups of URLs.

Official reference: Tavily API credits documentation.

2. Tavily plans and break-even analysis

Plan	Monthly price	Included credits	Effective rate
Researcher / Free	$0	1,000	Free within quota
Pay As You Go	Usage-based	No fixed bundle	$0.008 / credit
Project	$30	4,000	$0.0075 / included credit

Plan	Monthly price	Included credits	Effective rate
Bootstrap	$100	15,000	Approx. $0.00667
Startup	$220	38,000	Approx. $0.00579
Growth	$500	100,000	$0.005

Pay As You Go vs Project: $30 ÷ $0.008 = 3,750 credits. Below 3,750 credits, Pay As You Go is cheaper on the stated rates. At 4,000 credits, the $30 Project plan is cheaper than buying 4,000 credits individually ($32). The Project plan may also offer higher rate limits.

Overage: The earlier estimate assumes that usage beyond included credits can be billed at the stated pay-as-you-go rate when overage is enabled. Confirm the exact overage behavior and settings in your Tavily dashboard before relying on it.

Official reference: Tavily pricing.

3. LLM provider and model comparison

The following model prices were listed in the prior report as illustrative standard API rates per one million tokens. They are reproduced here for comparison, but must be rechecked against current provider documentation. Model identifiers and availability are particularly subject to change.

Provider	Model listed	Input / 1M	Output / 1M
Google	Gemini 2.5 Flash-Lite	$0.10	$0.40
Google	Gemini 2.5 Flash	$0.30	$2.50
Groq	GPT-OSS 20B	$0.075	$0.30
Groq	GPT-OSS 120B	$0.15	$0.60
Mistral	Ministral 3B	$0.10	$0.10
Mistral	Ministral 8B	$0.15	$0.15
Mistral	Mistral Small 4	$0.15	$0.60
OpenAI	GPT-6 Luna (verify model ID)	$0.05	$0.25
Anthropic	Claude Haiku 5.5 (verify model ID)	$0.10	$0.50

Do not use this table as a live quote. In particular, verify the exact model IDs for OpenAI and Anthropic in their official model catalogues before integrating them. The table is included to preserve the comparison made in the chat, not to guarantee that every listed model is currently orderable at those prices.

Free-tier and self-service availability

Provider	Free API tier	Paid self-service
Google Gemini	Yes, eligible models with quotas	Yes; separate API billing
Groq	Yes, rate-limited	Yes; Developer / usage-based
Mistral	Limited free access; varies by offer	Yes
OpenAI	No general free API allowance assumed	Yes; usage-based
Anthropic	No general free API allowance assumed	Yes; usage-based

Official references: Gemini · Groq · Mistral · OpenAI · Anthropic.

4. Google AI Pro subscription vs Gemini API

Google AI Pro is a consumer subscription. It does not provide a corresponding pool of Gemini API tokens or API credits. Gemini API usage is billed separately under the AI Studio / Gemini API billing arrangement.

Product	What it covers
Google AI Pro	Consumer Gemini app benefits and subscription usage limits.
Gemini API Free Tier	Free API usage for eligible models within published quotas.
Gemini API Paid Tier	Metered API usage under a separately configured billing account.

Use the existing Google AI Pro subscription for the consumer Gemini app, but do not budget it as free API capacity for your chatbot. For testing, the separate API free tier may be available. For identifiable resumes and personal career information, carefully review the applicable data-use terms; free and paid API tiers may have different data handling terms.

Official reference: Gemini API billing documentation.

5. Monthly traffic and token assumptions

Parameter	Low-traffic month	High-traffic month
Requests per day	100	1,000
Days per month	30	30
Total requests	3,000	30,000
Share requiring web search	20%	20%
Search calls / month	600	6,000
Input tokens per request	2,000	2,000
Output tokens per request	500	500
Total input tokens	6 million	60 million
Total output tokens	1.5 million	15 million

These are planning assumptions. The real token count depends on system prompts, conversation history, resume contents, retrieved pages, tool calls and model-specific tokenization.

6. Estimated LLM costs

Model	3,000 requests / month	30,000 requests / month
Gemini 2.5 Flash-Lite	$1.20	$12.00
Gemini 2.5 Flash	$5.55	$55.50
Groq GPT-OSS 20B	$0.90	$9.00
Groq GPT-OSS 120B	$1.80	$18.00
Mistral Ministral 3B	$0.75	$7.50
Mistral Ministral 8B	$1.13	$11.25
Mistral Small 4	$1.80	$18.00
OpenAI GPT-6 Luna (verify)	$0.68	$6.75
Claude Haiku 5.5 (verify)	$1.35	$13.50

Calculation assumes one generation per user request, 2,000 input tokens and 500 output tokens per request. Costs exclude cached-token discounts, reasoning-token surcharges, multimodal input, retries and any additional model calls.

7. Estimated Tavily costs

Scenario A — one Basic Search for every search-required request

Traffic	Credits needed	Free plan	Pure Pay As You Go	Project + assumed overage
Low	600	$0	$0 (within free quota)	Not needed
High	6,000	1,000 free; 5,000 paid	$48 gross; $40 after 1,000 free*	$46*

Scenario B — one Basic Search plus extraction of five pages for every search-required request

Traffic	Credits needed	Free plan	Pure Pay As You Go	Project + assumed overage
Low	1,200	1,000 free; 200 paid	$1.60 after free quota	Not needed
High	12,000	1,000 free; 11,000 paid	$96 gross; $88 after 1,000 free*	$94*

*The Project + overage estimates assume the $30 Project plan includes 4,000 credits, and additional credits cost $0.008 each. For 6,000 credits: $30 + 2,000 × $0.008 = $46. For 12,000 credits: $30 + 8,000 × $0.008 = $94. Confirm whether the free allowance and paid-plan credits stack in your account; do not assume that they do.

8. Combined estimated monthly costs

The following combines Gemini 2.5 Flash-Lite API usage with Tavily. It does not include Vercel hosting, Upstash Redis, embeddings, storage, taxes or payment fees.

Scenario	Low traffic	High traffic
Gemini 2.5 Flash-Lite	$1.20	$12.00
Tavily — search only	$0	$46.00*
Total — search only	$1.20	$58.00*
Tavily — search + extraction	$1.60	$94.00*
Total — search + extraction	$2.80	$106.00*

*High-traffic totals use the Project plan plus assumed overage, as calculated above. Actual costs depend on current plan rules and usage.

9. Recommended starting configuration

Component	Suggested starting choice	Reason
Search	Tavily Researcher (free)	Simple integration; measure real credit consumption first.
LLM	Gemini Flash-Lite API	Low estimated token cost; test answer quality.
Higher-quality model	Gemini Flash API	Use for complex resume analysis or planning if needed.
Alternative inference	Groq GPT-OSS 20B	Benchmark speed, quality and paid usage.
Cache	Upstash Redis	Cache reusable search results and safe general answers.

Upgrade guidance: During development, use free tiers with non-sensitive test data. At low usage, stay on Tavily's free plan if actual credit consumption fits. At sustained usage above roughly 3,750 credits per month, compare the $30 Project plan with Pay As You Go. Use the paid LLM tier and review data-use terms before handling real users' resumes.

10. Official pricing links

Tavily pricing · Tavily credit documentation

Google Gemini API pricing · Gemini API billing

Groq models · Groq billing

Mistral API pricing

OpenAI API pricing

Anthropic API pricing

Prepared as a planning report. Always verify current model availability, API prices, quotas, billing rules and data-use terms directly with providers before production deployment.